

Navyra Traffic Profile

2,000 prompts analysed · match radius 6

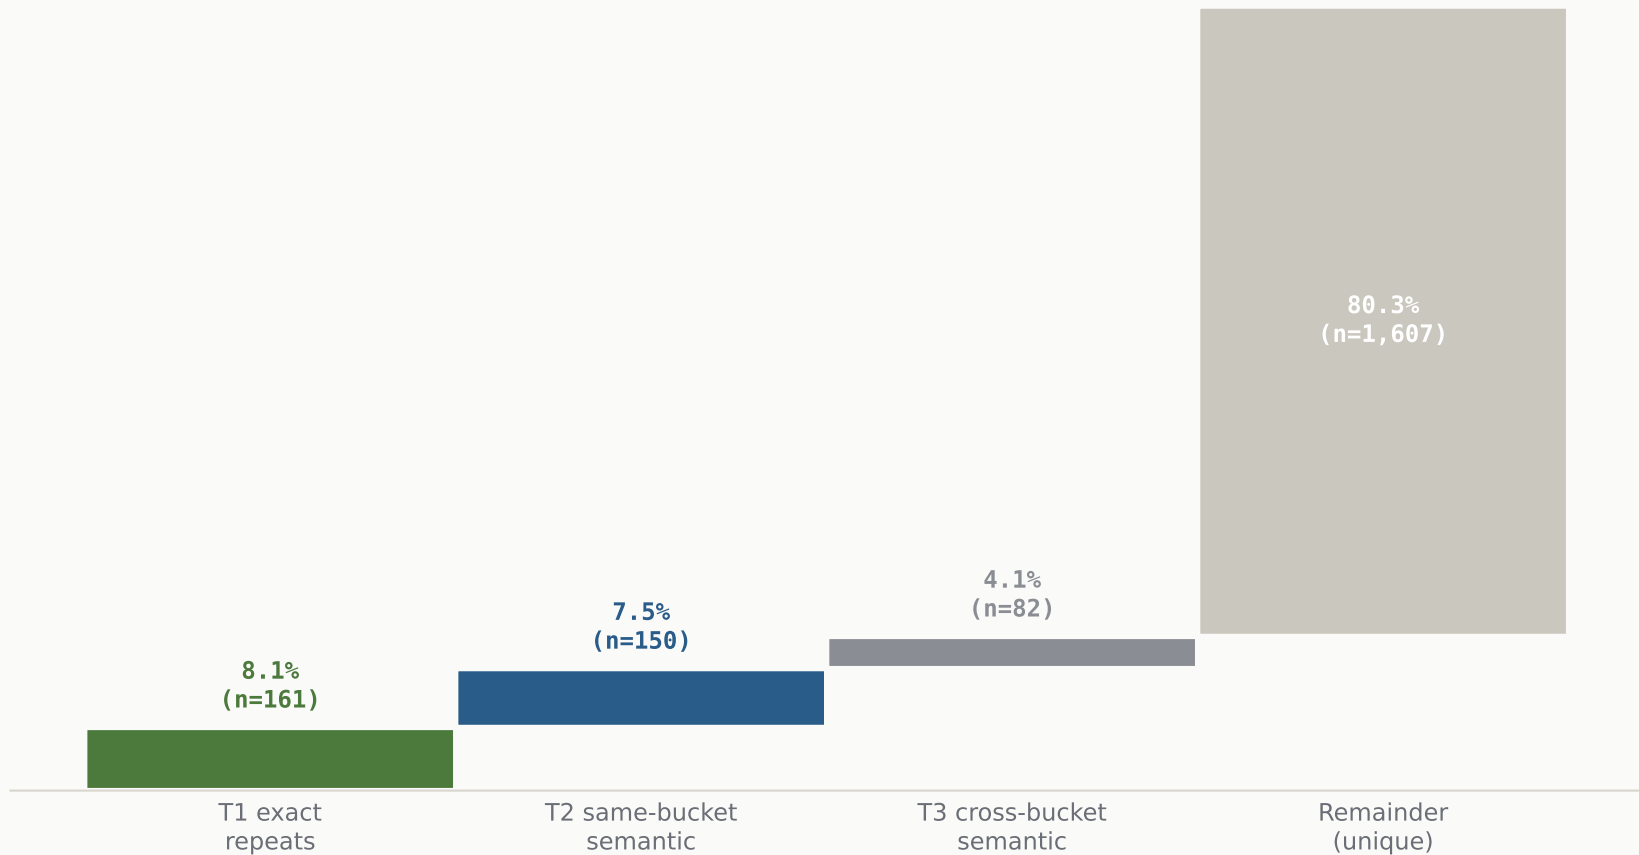
Note: T2 (same-bucket semantic) is an upper-bound proxy — up to this share of traffic is engine-addressable today. It is not a promised saving; the true hit rate is gated inside the model and measured in the free trial.

prompts analysed : 2,000
T1 exact repeats : 8.1% (already free via prefix caching)
T2 same-bucket semantic : 8.2% (engine-addressable today; upper-bound proxy)
T3 cross-bucket semantic : 4.5% (future-addressable)
T4 near-duplicates flagged: 5 (unsafe for answer-replay caches)
semantic clusters : 1,689
traffic in top cluster : 0.4%
hit rate by length bucket:
 bucket 8 tokens : 2.8% (n=977)
 bucket 16 tokens : 11.9% (n=1,002)
 bucket 24 tokens : 19.0% (n=21)

Based on 2,000 prompts · match radius 6

Tiered waterfall: where your traffic actually goes

Share of total traffic



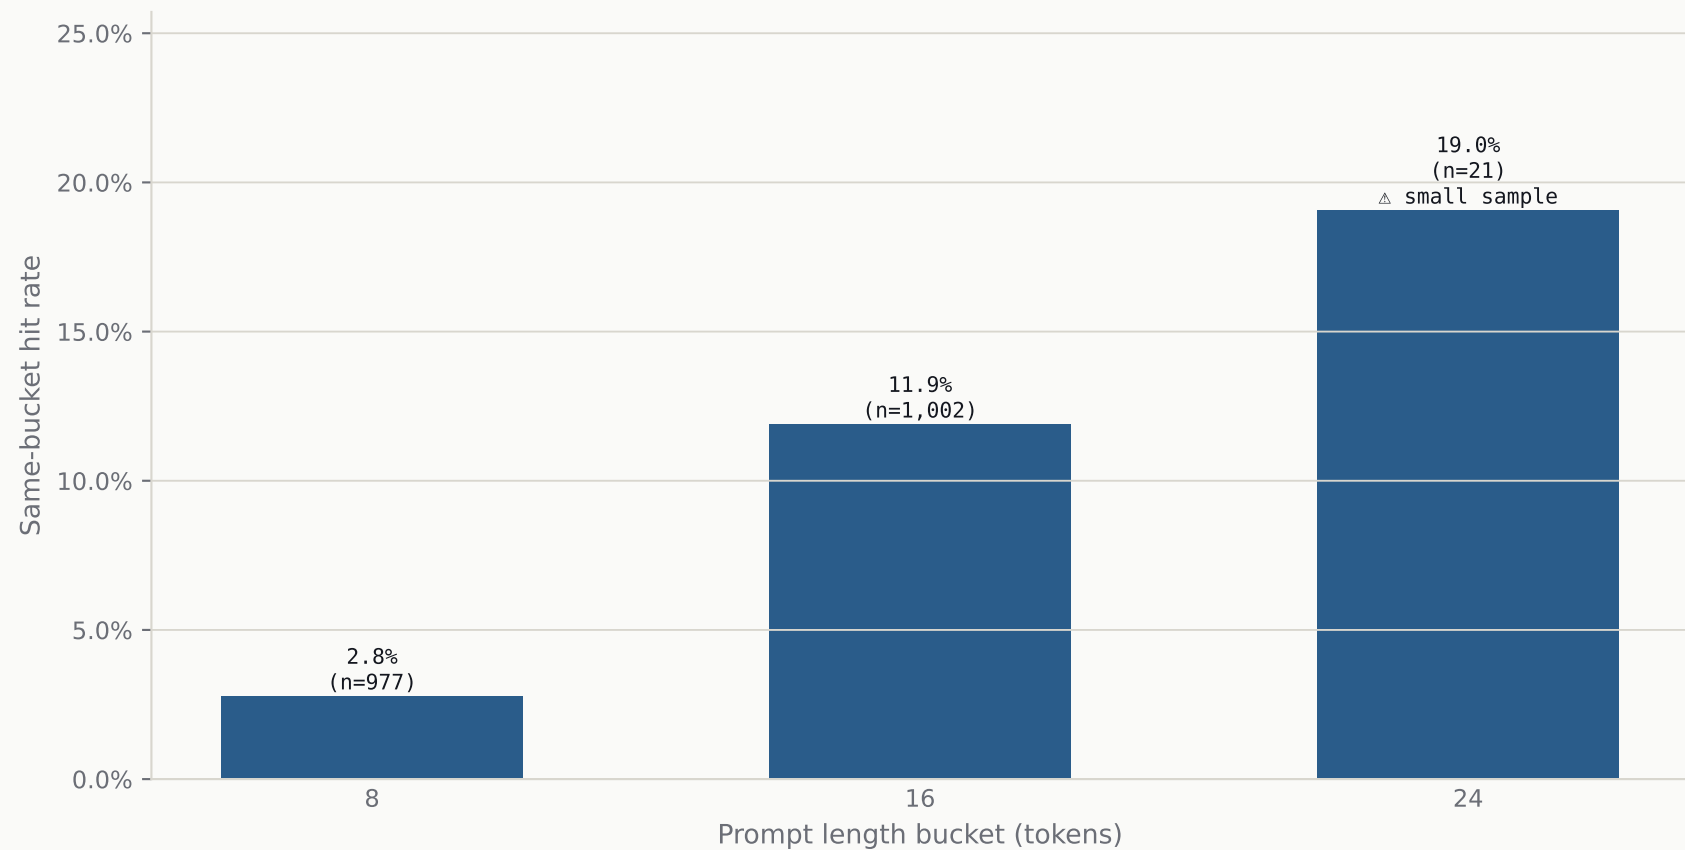
T1 — identical prompts, already served free by prefix caching.

T2 — different wording, same meaning, same length range: today's addressable opportunity (upper-bound estimate, not a guaranteed saving).

T3 — same meaning, different length range: not addressable yet.

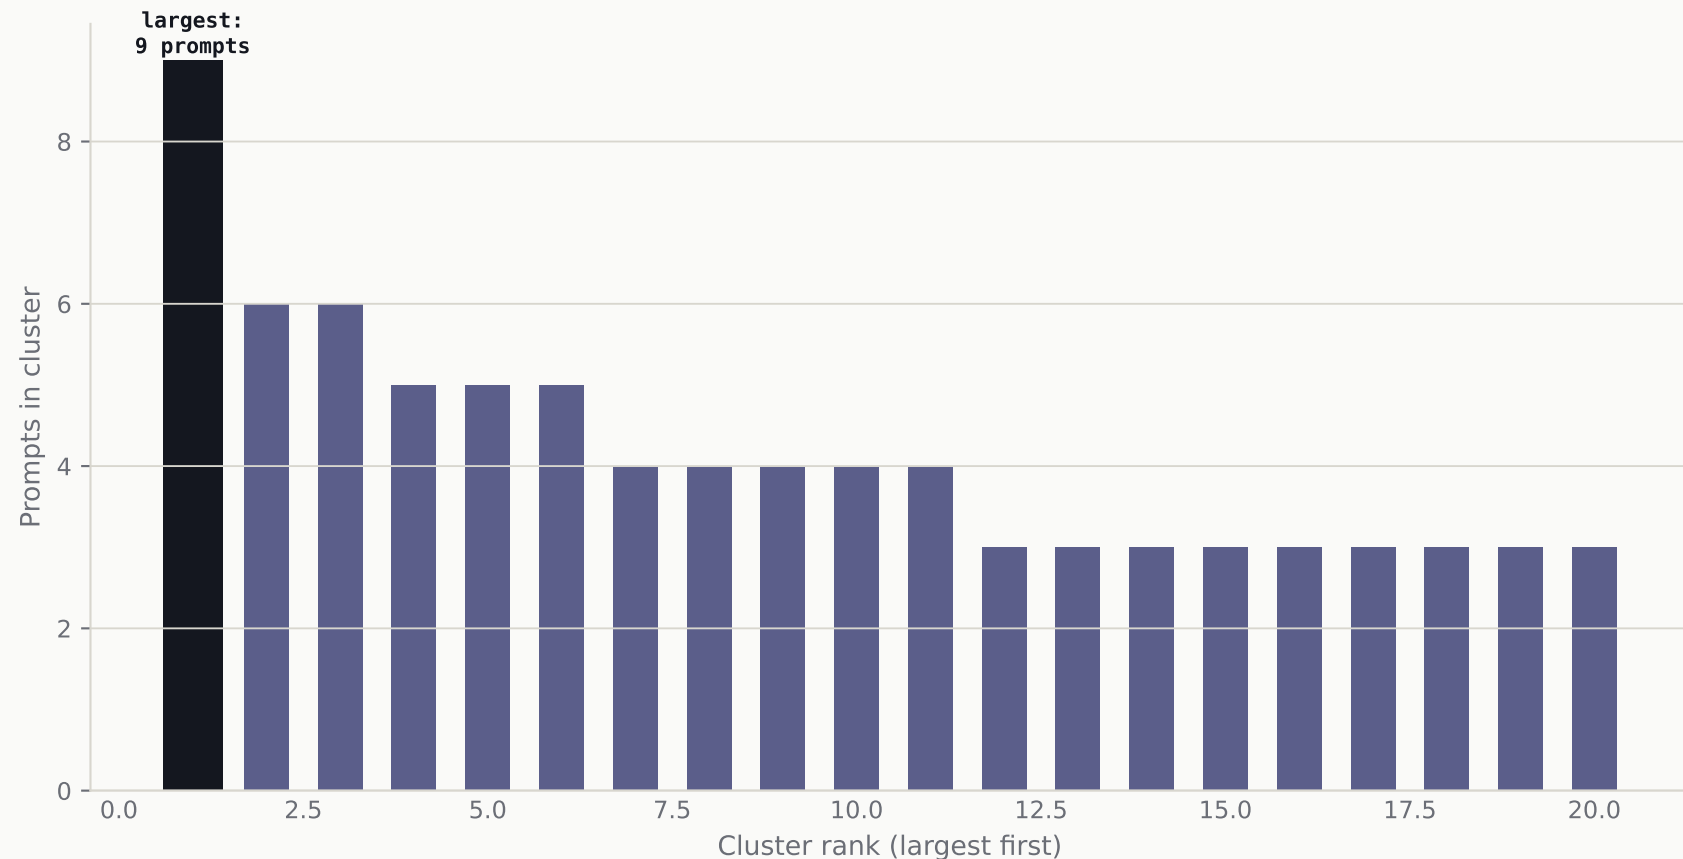
Remainder — traffic with no detected repetition.

Hit rate by length bucket



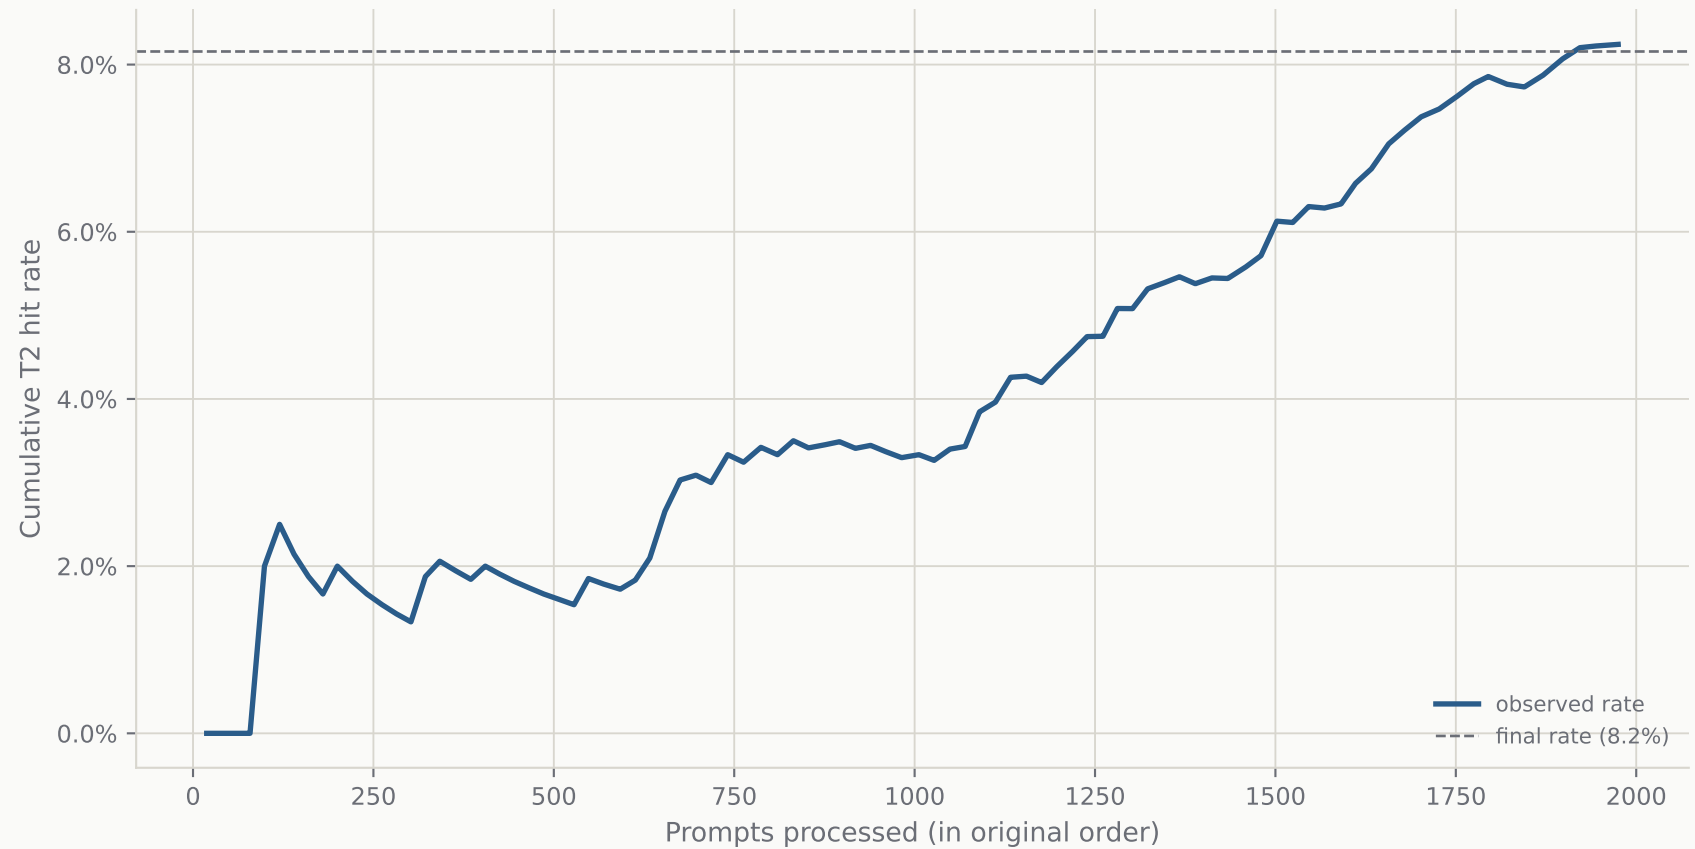
Each bar shows what share of prompts in that length range found a semantic match (T2) elsewhere in the same range. 'n' is the number of prompts that length range actually contains — small n means the rate is based on limited data and may not generalise.

Cluster size distribution — top 20 of 1,689 clusters total



A cluster is a group of prompts that all matched each other semantically. The 20 largest clusters shown here account for 33.6% of all traffic that formed any cluster (83 of 247 prompts). The single largest cluster alone is 0.4% of your total traffic.

Cache warm-up curve



This shows how the semantic hit rate changes as traffic streams in, simulating a cache starting empty. The rate is low early on (nothing to match against yet) and rises as more prompts accumulate in the bank. The dashed line marks the final overall rate for reference.